

Class 17: Vaccination Rate Mini Project

Suzanne Enos

Background

We're exploring a dataset on statewide vaccination rate from CA.gov

The goal of this hands-on mini-project is to examine and compare the Covid-19 vaccination rates around San Diego.

We will start by downloading the most recently dated "Statewide COVID-19 Vaccines Administered by ZIP Code" CSV file from: <https://data.ca.gov/dataset/covid-19-vaccine-progress-dashboard-data-by-zip-code>

Data Input

```
vax <- read.csv("covid19vaccinesbyzipcode_test.csv")
head(vax)
```

	as_of_date	zip_code_tabulation_area	local_health_jurisdiction	county
1	2021-01-05	93609	Fresno	Fresno
2	2021-01-05	94086	Santa Clara	Santa Clara
3	2021-01-05	94304	Santa Clara	Santa Clara
4	2021-01-05	94110	San Francisco	San Francisco
5	2021-01-05	93420	San Luis Obispo	San Luis Obispo
6	2021-01-05	93454	Santa Barbara	Santa Barbara
	vaccine_equity_metric_quartile	vem_source		
1		1	Healthy Places Index Score	
2		4	Healthy Places Index Score	
3		4	Healthy Places Index Score	
4		4	Healthy Places Index Score	
5		3	Healthy Places Index Score	
6		2	Healthy Places Index Score	

	age12_plus_population	age5_plus_population	tot_population
1	4396.3	4839	5177
2	42696.0	46412	50477
3	3263.5	3576	3852
4	64350.7	68320	72380
5	26694.9	29253	30740
6	32043.4	36446	40432

	persons_fully_vaccinated	persons_partially_vaccinated
1	NA	NA
2	11	640
3	NA	NA
4	18	1262
5	NA	NA
6	NA	NA

	percent_of_population_fully_vaccinated
1	NA
2	0.000218
3	NA
4	0.000249
5	NA
6	NA

	percent_of_population_partially_vaccinated
1	NA
2	0.012679
3	NA
4	0.017436
5	NA
6	NA

	percent_of_population_with_1_plus_dose	booster_recip_count
1	NA	NA
2	0.012897	NA
3	NA	NA
4	0.017685	NA
5	NA	NA
6	NA	NA

	bivalent_dose_recip_count	eligible_recipient_count
1	NA	1
2	NA	11
3	NA	6
4	NA	18
5	NA	4
6	NA	5

redacted

1 Information redacted in accordance with CA state privacy requirements
 2 Information redacted in accordance with CA state privacy requirements
 3 Information redacted in accordance with CA state privacy requirements
 4 Information redacted in accordance with CA state privacy requirements
 5 Information redacted in accordance with CA state privacy requirements
 6 Information redacted in accordance with CA state privacy requirements

Correct way to store numeric dates: year-month-day

Q1. What column details the total number of people fully vaccinated?

persons_fully_vaccinated

Q2. What column details the Zip code tabulation area?

zip_code_tabulation_area

What is the earliest date in this dataset?

2021-01-05

Q4. What is the latest date in this dataset?

```
tail(vax)
```

	as_of_date	zip_code_tabulation_area	local_health_jurisdiction
201091	2023-03-07	93662	Fresno
201092	2023-03-07	94801	Contra Costa
201093	2023-03-07	93668	Fresno
201094	2023-03-07	93704	Fresno
201095	2023-03-07	94510	Solano
201096	2023-03-07	93726	Fresno

	county	vaccine_equity_metric_quartile	vem_source
201091	Fresno	1	Healthy Places Index Score
201092	Contra Costa	1	Healthy Places Index Score
201093	Fresno	1	CDPH-Derived ZCTA Score
201094	Fresno	1	Healthy Places Index Score
201095	Solano	4	Healthy Places Index Score
201096	Fresno	1	Healthy Places Index Score

	age12_plus_population	age5_plus_population	tot_population
201091	24501.3	28311	30725
201092	25273.6	29040	31210
201093	1013.4	1199	1219

201094	24803.5	27701	29740
201095	24819.2	27056	28350
201096	33707.7	39067	42824
	persons_fully_vaccinated	persons_partially_vaccinated	
201091	20088	2150	
201092	27375	2309	
201093	644	74	
201094	17887	1735	
201095	22648	2264	
201096	24121	2682	
	percent_of_population_fully_vaccinated		
201091	0.653800		
201092	0.877123		
201093	0.528302		
201094	0.601446		
201095	0.798871		
201096	0.563259		
	percent_of_population_partially_vaccinated		
201091	0.069976		
201092	0.073983		
201093	0.060705		
201094	0.058339		
201095	0.079859		
201096	0.062628		
	percent_of_population_with_1_plus_dose	booster_recip_count	
201091	0.723776	10072	
201092	0.951106	14782	
201093	0.589007	312	
201094	0.659785	10435	
201095	0.878730	16092	
201096	0.625887	12104	
	bivalent_dose_recip_count	eligible_recipient_count	redacted
201091	2578	20066	No
201092	5342	27282	No
201093	66	644	No
201094	4154	17822	No
201095	8797	22501	No
201096	3585	24062	No

```
vax$as_of_date[nrow(vax)]
```

```
[1] "2023-03-07"
```

Use Skimr to get an overview of the dataset

If you just want to use one function from a package and not load the whole package: `package_name::` instead of `library()`

```
skimr::skim(vax)
```

Table 1: Data summary

Name	vax
Number of rows	201096
Number of columns	18
Column type frequency:	
character	5
numeric	13
Group variables	None

Variable type: character

skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
as_of_date	0	1	10	10	0	114	0
local_health_jurisdiction	0	1	0	15	570	62	0
county	0	1	0	15	570	59	0
vem_source	0	1	15	26	0	3	0
redacted	0	1	2	69	0	2	0

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
zip_code_tabulation_area	0	1.00	93665.11	1817.38	9000	192257.75	3658.50	5380.50	7635.0	
vaccine_equity_metric_99tile	0	0.95	2.44	1.11	1	1.00	2.00	3.00	4.0	
age12_plus_population	0	1.00	18895.04	8993.87	0	1346.95	13685.13	1756.18	8556.7	
age5_plus_population	0	1.00	20875.24	1105.97	0	1460.50	15364.06	14877.00	1902.0	
tot_population	9804	0.95	23372.77	2628.50	12	2126.00	18714.06	168.00	1165.0	

skim_variable	n_missing	complete	mean	sd	p0	p25	p50	p75	p100	hist
persons_fully_vaccinated	16621	0.92	13990.30	5073.66	11	932.00	8589.00	23346.00	87575.0	
persons_partially_vaccinated	16621	0.92	1702.31	2033.32	11	165.00	1197.00	2536.00	39973.0	
percent_of_population_fully_vaccinated	20965	0.90	0.57	0.25	0	0.42	0.61	0.74	1.0	
percent_of_population_partially_vaccinated	20965	0.90	0.08	0.09	0	0.05	0.06	0.08	1.0	
percent_of_population_1_plus_dose	20965	0.89	0.63	0.24	0	0.49	0.67	0.81	1.0	
booster_recip_count	72997	0.64	5882.76	219.00	11	300.00	2773.00	9510.00	59593.0	
bivalent_dose_recip_count	158776	0.21	2978.23	3633.03	11	193.00	1467.50	4730.25	27694.0	
eligible_recipient_count	0	1.00	12830.83	4928.64	0	507.00	6369.00	2014.00	87248.0	

Q5. How many numeric columns are in this dataset?

13

Q6. Note that there are “missing values” in the dataset. How many NA values there in the persons_fully_vaccinated column?

16621

```
n.missing <- sum(is.na(vax$persons_fully_vaccinated))
n.missing
```

[1] 16621

Q7. What percent of persons_fully_vaccinated values are missing (to 2 significant figures)?

```
round(n.missing/nrow(vax)*100, 2)
```

[1] 8.27

Working with dates

We will use the lubridate package to help ease the pain of working with times and dates

```
library(lubridate)
```

Attaching package: 'lubridate'

The following objects are masked from 'package:base':

date, intersect, setdiff, union

```
today()
```

```
[1] "2023-03-08"
```

```
today() - ymd(vax$as_of_date[1])
```

Time difference of 792 days

How old am I in days?

```
today() - ymd("2000-01-27")
```

Time difference of 8441 days

Q9. How many days have passed since the last update of the dataset?

```
today() - ymd(vax$as_of_date[nrow(vax)])
```

Time difference of 1 days

```
vax$as_of_date <- ymd(vax$as_of_date)
```

```
today() - vax$as_of_date[nrow(vax)]
```

Time difference of 1 days

Q10. How many unique dates are in the dataset (i.e. how many different dates are detailed)?

```
length(unique(vax$as_of_date))
```

```
[1] 114
```

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

```
filter, lag
```

The following objects are masked from 'package:base':

```
intersect, setdiff, setequal, union
```

```
n_distinct(vax$as_of_date)
```

```
[1] 114
```

Working with zipcodes

Can use zipcodeR package to make this easier to work with

```
library(zipcodeR)
geocode_zip("30022")
```

```
# A tibble: 1 x 3
  zipcode lat lng
  <chr>   <dbl> <dbl>
1 30022   34.0 -84.2
```

```
zip_distance('92037', '92109')
```

```
zipcode_a zipcode_b distance
1      92037      92109      2.33
```

```
reverse_zipcode(c('92037', '92109', '30605'))
```



```
# A tibble: 3 x 24
  zipcode zipcode~1 major~2 post_~3 common_c~4 county state lat lng timez~5
  <chr> <chr> <chr> <chr> <blob> <chr> <chr> <dbl> <dbl> <chr>
1 92037 Standard La Jol~ La Jol~ <raw 20 B> San D~ CA 32.8 -117. Pacific
2 92109 Standard San Di~ San Di~ <raw 21 B> San D~ CA 32.8 -117. Pacific
3 30605 Standard Athens Athens~ <raw 18 B> Clark~ GA 33.9 -83.3 Eastern
# ... with 14 more variables: radius_in_miles <dbl>, area_code_list <blob>,
# population <int>, population_density <dbl>, land_area_in_sqmi <dbl>,
# water_area_in_sqmi <dbl>, housing_units <int>,
# occupied_housing_units <int>, median_home_value <int>,
# median_household_income <int>, bounds_west <dbl>, bounds_east <dbl>,
# bounds_north <dbl>, bounds_south <dbl>, and abbreviated variable names
# 1: zipcode_type, 2: major_city, 3: post_office_city, ...
```

Finding best and worst ratio of household income to home prices

Find all SD zipcodes

```
sd.zip <- unique(vax$zip_code_tabulation_area[vax$county == "San Diego"])
length(sd.zip)
```

[1] 107

```
sd.eco <- reverse_zipcode(sd.zip)
sd.eco
```

```
# A tibble: 107 x 24
  zipcode zipcode~1 major~2 post_~3 common_c~4 county state lat lng timez~5
  <chr> <chr> <chr> <chr> <blob> <chr> <chr> <dbl> <dbl> <chr>
1 91901 Standard Alpine Alpine~ <raw 18 B> San D~ CA 32.8 -117. Pacific
2 91902 Standard Bonita Bonita~ <raw 18 B> San D~ CA 32.7 -117. Pacific
3 91905 Standard Boulev~ Boulev~ <raw 21 B> San D~ CA 32.7 -116. Pacific
4 91906 Standard Campo Campo,~ <raw 17 B> San D~ CA 32.7 -116. Pacific
5 91910 Standard Chula ~ Chula ~ <raw 23 B> San D~ CA 32.6 -117. Pacific
6 91911 Standard Chula ~ Chula ~ <raw 23 B> San D~ CA 32.6 -117. Pacific
7 91913 Standard Chula ~ Chula ~ <raw 23 B> San D~ CA 32.6 -117. Pacific
8 91914 Standard Chula ~ Chula ~ <raw 23 B> San D~ CA 32.7 -117. Pacific
9 91915 Standard Chula ~ Chula ~ <raw 23 B> San D~ CA 32.6 -117. Pacific
10 91916 Standard Descan~ Descan~ <raw 20 B> San D~ CA 32.9 -117. Pacific
# ... with 97 more rows, 14 more variables: radius_in_miles <dbl>,
# area_code_list <blob>, population <int>, population_density <dbl>,
```

```
# land_area_in_sqmi <dbl>, water_area_in_sqmi <dbl>, housing_units <int>,
# occupied_housing_units <int>, median_home_value <int>,
# median_household_income <int>, bounds_west <dbl>, bounds_east <dbl>,
# bounds_north <dbl>, bounds_south <dbl>, and abbreviated variable names
# 1: zipcode_type, 2: major_city, 3: post_office_city, ...
```

Most expensive area?

```
ord <- order(sd.eco$median_household_income, decreasing=T)
head(sd.eco[ord,])
```

```
# A tibble: 6 x 24
  zipcode zipcode_~1 major~2 post_~3 common_c~4 county state lat lng timez~5
  <chr> <chr> <chr> <chr> <blob> <chr> <chr> <dbl> <dbl> <chr>
1 92145 Unique San Di~ San Di~ <raw 21 B> San D~ CA 32.9 -117. Pacific
2 92091 Standard Rancho~ Rancho~ <raw 33 B> San D~ CA 33 -117. Pacific
3 92130 Standard San Di~ San Di~ <raw 21 B> San D~ CA 33.0 -117. Pacific
4 92067 PO Box Rancho~ Rancho~ <raw 33 B> San D~ CA 33.0 -117. Pacific
5 92131 Standard San Di~ San Di~ <raw 21 B> San D~ CA 32.9 -117. Pacific
6 92127 Standard San Di~ San Di~ <raw 21 B> San D~ CA 33.0 -117. Pacific
# ... with 14 more variables: radius_in_miles <dbl>, area_code_list <blob>,
# population <int>, population_density <dbl>, land_area_in_sqmi <dbl>,
# water_area_in_sqmi <dbl>, housing_units <int>,
# occupied_housing_units <int>, median_home_value <int>,
# median_household_income <int>, bounds_west <dbl>, bounds_east <dbl>,
# bounds_north <dbl>, bounds_south <dbl>, and abbreviated variable names
# 1: zipcode_type, 2: major_city, 3: post_office_city, ...
```

```
head(arrange(sd.eco, desc(median_home_value)))
```

```
# A tibble: 6 x 24
  zipcode zipcode_~1 major~2 post_~3 common_c~4 county state lat lng timez~5
  <chr> <chr> <chr> <chr> <blob> <chr> <chr> <dbl> <dbl> <chr>
1 92014 Standard Del Mar Del Ma~ <raw 19 B> San D~ CA 33.0 -117. Pacific
2 92037 Standard La Jol~ La Jol~ <raw 20 B> San D~ CA 32.8 -117. Pacific
3 92067 PO Box Rancho~ Rancho~ <raw 33 B> San D~ CA 33.0 -117. Pacific
4 92118 Standard Corona~ Corona~ <raw 33 B> San D~ CA 32.6 -117. Pacific
5 92145 Unique San Di~ San Di~ <raw 21 B> San D~ CA 32.9 -117. Pacific
6 92091 Standard Rancho~ Rancho~ <raw 33 B> San D~ CA 33 -117. Pacific
# ... with 14 more variables: radius_in_miles <dbl>, area_code_list <blob>,
```

```
# population <int>, population_density <dbl>, land_area_in_sqmi <dbl>,
# water_area_in_sqmi <dbl>, housing_units <int>,
# occupied_housing_units <int>, median_home_value <int>,
# median_household_income <int>, bounds_west <dbl>, bounds_east <dbl>,
# bounds_north <dbl>, bounds_south <dbl>, and abbreviated variable names
# 1: zipcode_type, 2: major_city, 3: post_office_city, ...
```

```
sd.eco$median_home_value
```

```
[1] 472500 495600 273700 187100 340300 294900 387800 461300 365800
[10] 254700 182800 70700 340900 145600 503900 489200 344600 287700
[19]      NA 263500 336200 78200 296400 344400      NA 521500 196300
[28] 743200 615200 641100 484200 665300 1000001 405000 380200 299000
[37] 685700 404000 359500 304100 423600 486900 320500 1000001 340900
[46] 426200      NA 353100 341000 261900 298700      NA 356100 495800
[55] 383500 300000 1000001 362800 362000 323500 950300 369700 390700
[64] 439800 284000 391800 169000 986000 448700 278500 546800 427200
[73] 266700 834800 650700 285000 717300 466300 387100 238200 272300
[82] 365800 447200 447300 1000001 423700 455800 569000 576200 396100
[91] 519100 378400 618400 491300 577700 832900 627000      NA      NA
[100]      NA 273400      NA 1000001      NA 269800      NA 254300
```

Focus on the San Diego area

Using `dplyr` here

```
sd <- filter(vax, county == "San Diego")

nrow(sd)
```

```
[1] 12198
```

Q11. How many distinct zip codes are listed for San Diego County?

```
n_distinct(sd$zip_code_tabulation_area)
```

```
[1] 107
```

Q12. What San Diego County Zip code area has the largest 12 + Population in this dataset?

```
ind <- which.max(sd$age12_plus_population)
sd[ind,]
```

```

as_of_date zip_code_tabulation_area local_health_jurisdiction county
67 2021-01-05 92154 San Diego San Diego
vaccine_equity_metric_quartile vem_source
67 2 Healthy Places Index Score
age12_plus_population age5_plus_population tot_population
67 76365.2 82971 88979
persons_fully_vaccinated persons_partially_vaccinated
67 16 1400
percent_of_population_fully_vaccinated
67 0.00018
percent_of_population_partially_vaccinated
67 0.015734
percent_of_population_with_1_plus_dose booster_recip_count
67 0.015914 NA
bivalent_dose_recip_count eligible_recipient_count
67 NA 16
redacted
67 Information redacted in accordance with CA state privacy requirements

```

Q13. What is the overall average “Percent of Population Fully Vaccinated” value for all San Diego “County” as of “2023-03-07”?

```

library(dplyr)
sd.latest <- filter(sd, as_of_date == "2023-03-07")
mean(sd.latest$percent_of_population_fully_vaccinated, na.rm = TRUE)

```

```
[1] 0.7402567
```

Q14. Using either ggplot or base R graphics make a summary figure that shows the distribution of Percent of Population Fully Vaccinated values as of “2023-02-28”?


```

library(ggplot2)
ggplot(sd.latest) +
  aes(sd.latest$percent_of_population_fully_vaccinated) +
  geom_histogram(bins = 20) +
  theme_bw() +
  xlab("Percent of People Vaccinated")

```

Warning: Use of `sd.latest\$percent_of_population_fully_vaccinated` is discouraged.
i Use `percent_of_population_fully_vaccinated` instead.

Warning: Removed 8 rows containing non-finite values (`stat_bin()`).




Focus on UCSD/La Jolla

```
ucsd <- filter(sd, zip_code_tabulation_area == "92037")
ucsd[1,]$age5_plus_population
```

```
[1] 36144
```

```
ljplot <- ggplot(ucsd) +
  aes(x = as_of_date, y = percent_of_population_fully_vaccinated) +
  geom_point() +
  geom_line(group=1) +
  ylim(c(0,1)) +
  labs(title = "Vaccination rate for La Jolla CA 92037", x = "Date", y = "Percent Vaccinated")
```

```
theme_bw()
ljplot
```



Comparing to similar sized areas

```
vax.36 <- filter(vax, age5_plus_population > 36144 &
  as_of_date == "2023-03-07")
head(vax.36)
```

	as_of_date	zip_code_tabulation_area	local_health_jurisdiction	county
1	2023-03-07	94116	San Francisco	San Francisco
2	2023-03-07	92703	Orange	Orange
3	2023-03-07	94118	San Francisco	San Francisco
4	2023-03-07	92376	San Bernardino	San Bernardino
5	2023-03-07	92692	Orange	Orange
6	2023-03-07	95148	Santa Clara	Santa Clara

	vaccine_equity_metric_quartile	vem_source
1	4	Healthy Places Index Score
2	1	Healthy Places Index Score

3		4 Healthy Places Index Score	
4		1 Healthy Places Index Score	
5		4 Healthy Places Index Score	
6		4 Healthy Places Index Score	
	age12_plus_population	age5_plus_population	tot_population
1	42334.3	45160	47346
2	57182.7	64387	69112
3	37628.5	40012	42095
4	70232.1	79686	86085
5	41008.9	44243	46800
6	42163.3	46202	48273
	persons_fully_vaccinated	persons_partially_vaccinated	
1	41255	2450	
2	57887	7399	
3	33284	3040	
4	51367	5674	
5	35117	2603	
6	42298	2684	
	percent_of_population_fully_vaccinated		
1	0.871351		
2	0.837582		
3	0.790688		
4	0.596701		
5	0.750363		
6	0.876225		
	percent_of_population_partially_vaccinated		
1	0.051747		
2	0.107058		
3	0.072218		
4	0.065912		
5	0.055620		
6	0.055600		
	percent_of_population_with_1_plus_dose	booster_recip_count	
1	0.923098	34108	
2	0.944640	28297	
3	0.862906	27401	
4	0.662613	23832	
5	0.805983	23695	
6	0.931825	31583	
	bivalent_dose_recip_count	eligible_recipient_count	redacted
1	19158	41000	No
2	7627	57775	No
3	15251	33146	No

4	6393	51276	No
5	10169	35031	No
6	12604	42120	No

How many unique zipcodes?


```
n_distinct(vax.36$zip_code_tabulation_area)
```

[1] 411

Q16. Calculate the mean “Percent of Population Fully Vaccinated” for ZIP code areas with a population as large as 92037 (La Jolla) as_of_date “2023-02-28”. Add this as a straight horizontal line to your plot from above with the `geom_hline()` function

```
avg.36 <- mean(vax.36$percent_of_population_fully_vaccinated)
```

```
ljplot +  
  geom_hline(yintercept = avg.36, color = "red", linetype = 2)
```



Q17. What is the 6 number summary (Min, 1st Qu., Median, Mean, 3rd Qu., and

Max) of the “Percent of Population Fully Vaccinated” values for ZIP code areas with a population as large as 92037 (La Jolla) as_of_date “2023-02-28”?

```
summary(vax.36$percent_of_population_fully_vaccinated)
```

```
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.3805  0.6459  0.7183  0.7215  0.7908  1.0000
```


Q18. Using ggplot generate a histogram of this data.

```
ggplot(vax.36) +
  aes(vax.36$percent_of_population_fully_vaccinated) +
  geom_histogram() +
  xlim(0,1) +
  labs(x = "Percent Vaccinated")
```

Warning: Use of `vax.36\$percent_of_population_fully_vaccinated` is discouraged.
i Use `percent_of_population_fully_vaccinated` instead.

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

Warning: Removed 2 rows containing missing values (`geom_bar()`).



Q19. Is the 92109 and 92040 ZIP code areas above or below the average value you calculated for all these above?

```
filter(vax.36, zip_code_tabulation_area %in% c("92109", "92040"))
```

	as_of_date	zip_code_tabulation_area	local_health_jurisdiction	county
1	2023-03-07	92109	San Diego	San Diego
2	2023-03-07	92040	San Diego	San Diego
	vaccine_equity_metric_quartile	vem_source		
1	3	Healthy Places Index Score		
2	3	Healthy Places Index Score		
	age12_plus_population	age5_plus_population	tot_population	
1	43222.5	44953	47111	
2	39405.0	42833	46306	
	persons_fully_vaccinated	persons_partially_vaccinated		
1	32725	4234		
2	25493	2156		
	percent_of_population_fully_vaccinated			
1	0.694636			
2	0.550533			
	percent_of_population_partially_vaccinated			
1	0.089873			
2	0.046560			
	percent_of_population_with_1_plus_dose	booster_recip_count		
1	0.784509	19677		
2	0.597093	14175		
	bivalent_dose_recip_count	eligible_recipient_count	redacted	
1	8109	32622	No	
2	4649	25433	No	

Q20. Finally make a time course plot of vaccination progress for all areas in the full dataset with a age5_plus_population > 36144.

```
vax.36.all <- filter(vax, age5_plus_population > 36144)
```

```
ggplot(vax.36.all) +
  aes(x = as_of_date, y = percent_of_population_fully_vaccinated, group = zip_code_tabulation_area) +
  geom_line(alpha = 0.2, color = "blue") +
  ylim(0,1) +
  labs(x = "Date", y = "Percent Vaccinated", title = "Vaccination rate across California",
  geom_hline(yintercept = avg.36, linetype = 2)
```

Warning: Removed 183 rows containing missing values (``geom_line()``).

